

```
In [2]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn import preprocessing, svm
from sklearn.linear_model import LinearRegression
```

```
In [3]: df=pd.read_csv(r"C:\Users\Lenovo\OneDrive\Desktop\Data Sets\fiat500_VehicleSele
df
```

Out[3]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611560
1	2	pop	51	1186	32500	1	45.666359	12.241890
2	3	sport	74	4658	142228	1	45.503300	11.417840
3	4	lounge	51	2739	160000	1	40.633171	17.634609
4	5	pop	73	3074	106880	1	41.903221	12.495650
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870
1535	1536	pop	51	2223	60457	1	45.481541	9.413480
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270
1537	1538	pop	51	1766	54276	1	40.323410	17.568270

1538 rows × 9 columns



```
In [4]: df.head()
```

Out[4]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700



In [5]: `df.tail()`

Out[5]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
1533	1534	sport	51	3712	115280	1	45.069679	7.70492	153300000
1534	1535	lounge	74	3835	112000	1	45.845692	8.66687	153400000
1535	1536	pop	51	2223	60457	1	45.481541	9.41348	153500000
1536	1537	lounge	51	2557	80750	1	45.000702	7.68227	153600000
1537	1538	pop	51	1766	54276	1	40.323410	17.56827	153700000

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    1538 non-null   int64
1   model                 1538 non-null   object
2   engine_power          1538 non-null   int64
3   age_in_days           1538 non-null   int64
4   km                    1538 non-null   int64
5   previous_owners       1538 non-null   int64
6   lat                   1538 non-null   float64
7   lon                   1538 non-null   float64
8   price                 1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [8]: `df.describe()`

Out[8]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	12.133518
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.133518
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	12.133518
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	12.133518
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	12.133518
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.133518
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	12.133518

In [10]: `df.shape`

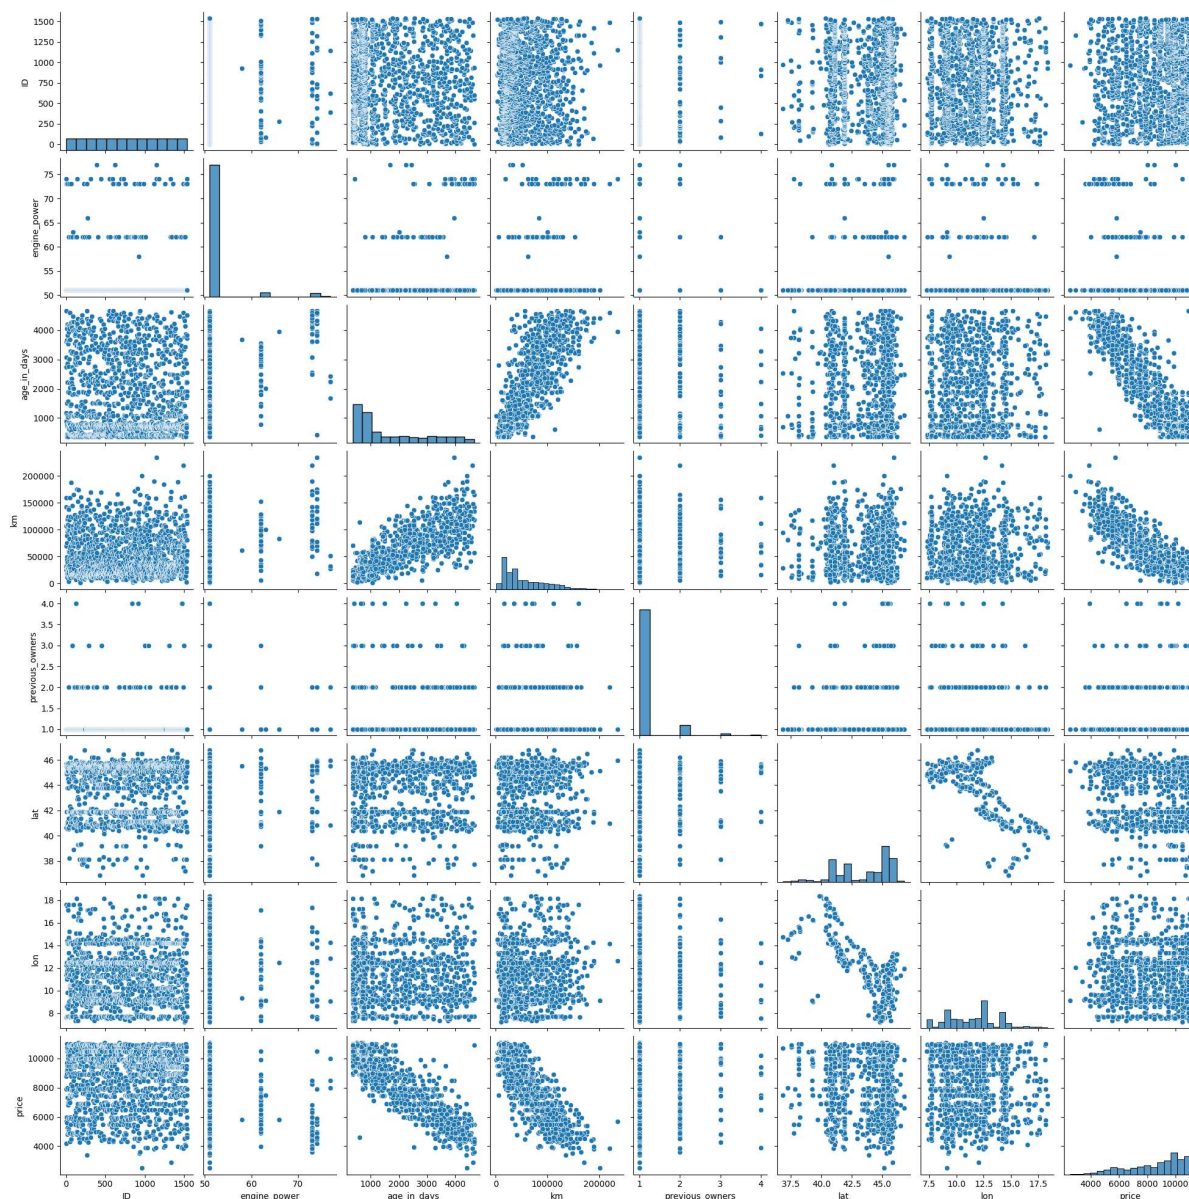
Out[10]: (1538, 9)

```
In [12]: df.columns
```

```
Out[12]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
              'lat', 'lon', 'price'],  
              dtype='object')
```


```
In [13]: sns.pairplot(df)
```

```
Out[13]: <seaborn.axisgrid.PairGrid at 0x2c2345047d0>
```




```
In [14]: sns.displot(['Price'])
```

```
Out[14]: <seaborn.axisgrid.FacetGrid at 0x2c24bce64d0>
```




```
In [15]: sns.displot(df['price'])
```

```
Out[15]: <seaborn.axisgrid.FacetGrid at 0x2c24bac3e10>
```




```
In [16]: sns.displot(df['engine_power'])
```

```
Out[16]: <seaborn.axisgrid.FacetGrid at 0x2c24bd28650>
```




```
In [17]: fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
                  'lat', 'lon','price']]  
sns.heatmap(fiatdf.corr())
```

Out[17]: <Axes: >



```
In [18]: fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',
'lat', 'lon']]
sns.heatmap(fiatdf.corr())#without price
```

Out[18]: <Axes: >



```
In [19]: X=fiatdf[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',
'lat', 'lon']]
y=df['price']
```

```
In [20]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=1)
from sklearn.linear_model import LinearRegression
regr=LinearRegression()
regr.fit(X_train,y_train)
print(regr.intercept_)
```

8971.195683500262



```
In [21]: coeff_df=pd.DataFrame(regr.coef_,X.columns,columns=['coefficient'])  
coeff_df
```

Out[21]:

	coefficient
ID	-0.046704
engine_power	11.646408
age_in_days	-0.898018
km	-0.017232
previous_owners	26.400886
lat	32.189709
lon	0.161073


```
In [22]: predictions=regr.predict(X_test)  
plt.scatter(y_test,predictions)
```

Out[22]: <matplotlib.collections.PathCollection at 0x2c24e17fa50>




```
In [23]: sns.displot((y_test,predictions),bins=50)#without semicolon
```

```
Out[23]: <seaborn.axisgrid.FacetGrid at 0x2c24de1c7d0>
```



```
In [24]: sns.displot((y_test,predictions),bins=50);#with semicolon
```



```
In [25]: from sklearn import metrics
print('MAE:',metrics.mean_absolute_error(y_test,predictions))
print('MSE:',metrics.mean_squared_error(y_test,predictions))
print('MAE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

MAE: 593.0876179519996

MSE: 551442.6799691911

MAE: 742.59186635001

```
In [26]: #accuracy
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

0.859713670430884


```
In [27]: df.fillna(method='ffill',inplace=True)
```

```
In [28]: x=np.array(df['age_in_days']).reshape(-1,1)
y=np.array(df['km']).reshape(-1,1)
df.dropna(inplace=True)
```

```
In [29]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
```

```
Out[29]: ▾ LinearRegression
LinearRegression()
```

```
In [30]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



```
In [32]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
y_pred_elastic=regr.predict(X_train)
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

[25.89689696]

[10640.73996329]

Mean Squared Error on test set 2740765507.785756

In []: